

Concepts

Core architecture

Understand how MCP connects clients, servers, and LLMs

The Model Context Protocol (MCP) is built on a flexible, extensible architecture that enables seamless communication between LLM applications and integrations. This document covers the core architectural components and concepts.


Overview

MCP follows a client-server architecture where:

Hosts are LLM applications (like Claude Desktop or IDEs) that initiate connections

Clients maintain 1:1 connections with servers, inside the host application

Servers provide context, tools, and prompts to clients



Protocol layer

Concepts > **Core architecture**

The protocol layer handles message framing, request/response linking, and high-level communication patterns.

TypeScript Python

```
class Protocol<Request, Notification, Result> {
  // Handle incoming requests
  setRequestHandler<T>(schema: T, handler: (request: T, extra: R) => Promise<Result>)

  // Handle incoming notifications
  setNotificationHandler<T>(schema: T, handler: (notification: T) => Promise<void>)

  // Send requests and await responses
  request<T>(request: Request, schema: T, options?: RequestOptions)

  // Send one-way notifications
  notification(notification: Notification): Promise<void>
}
```

Key classes include:

Protocol

Client

Server

Transport layer

The transport layer handles the actual communication between clients and servers. MCP supports multiple transport mechanisms:

1. Stdio transport



2. HTTP with SSE transport

Uses Server-Sent Events for server-to-client messages

HTTP POST for client-to-server messages

All transports use **JSON-RPC 2.0** to exchange messages. See the [specification](#) for detailed information about the Model Context Protocol message format.

Message types

MCP has these main types of messages:

1. **Requests** expect a response from the other side:

```
interface Request {  
  method: string;  
  params?: { ... };  
}
```

2. **Results** are successful responses to requests:

```
interface Result {  
  [key: string]: unknown;  
}
```

3. **Errors** indicate that a request failed:

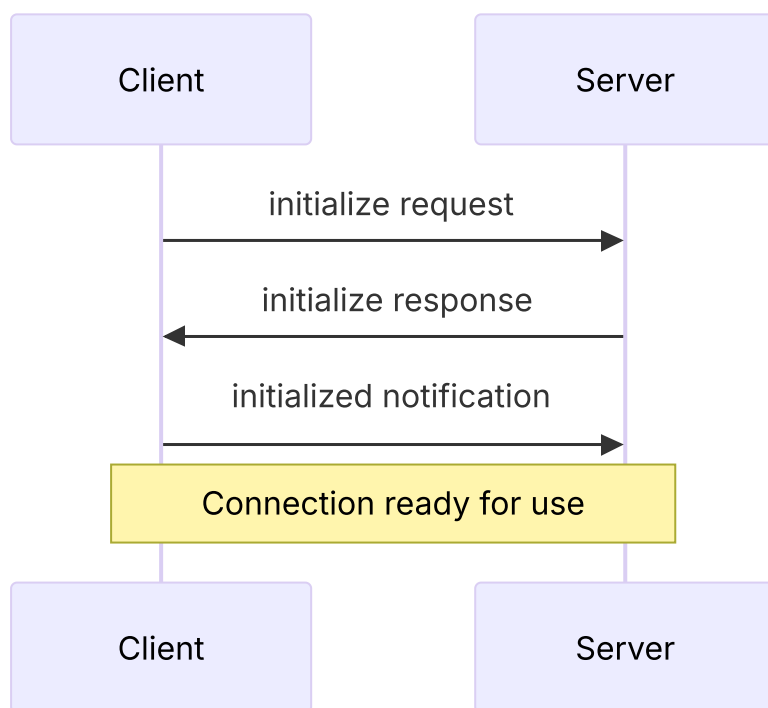
```
interface Error {  
  code: number;  
  message: string;
```

4. Notifications are one-way messages that don't expect a response:

```
interface Notification {  
  method: string;  
  params?: { ... };  
}
```

Connection lifecycle

1. Initialization



1. Client sends `initialize` request with protocol version and capabilities
2. Server responds with its protocol version and capabilities
3. Client sends `initialized` notification as acknowledgment
4. Normal message exchange begins

2. Message exchange



After initialization, the following patterns are supported:
Model Context Protocol

Request-Response: Client or server sends requests, the other responds

Concepts and architecture

Notifications: Either party sends one-way messages

3. Termination

Either party can terminate the connection:

- Clean shutdown via `close()`

- Transport disconnection

- Error conditions

Error handling

MCP defines these standard error codes:

```
enum ErrorCode {  
    // Standard JSON-RPC error codes  
    ParseError = -32700,  
    InvalidRequest = -32600,  
    MethodNotFound = -32601,  
    InvalidParams = -32602,  
    InternalError = -32603  
}
```

SDKs and applications can define their own error codes above -32000.

Errors are propagated through:

- Error responses to requests

- Error events on transports

- Protocol-level error handlers



Implementation example

Model Context Protocol

Here's a basic example of implementing an MCP server:

Concepts > **Core architecture**

TypeScript Python

```
import { Server } from "@modelcontextprotocol/sdk/server/index.js"
import { StdioServerTransport } from "@modelcontextprotocol/sdk/server/stdio.js"

const server = new Server({
  name: "example-server",
  version: "1.0.0"
}, {
  capabilities: {
    resources: {}
  }
});

// Handle requests
server.setRequestHandler(ListResourcesRequestSchema, async () => {
  return {
    resources: [
      {
        uri: "example://resource",
        name: "Example Resource"
      }
    ]
  };
});

// Connect transport
const transport = new StdioServerTransport();
await server.connect(transport);
```

Best practices

Transport selection



1. Local communication

Model Context Protocol

Use stdio transport for local processes

Concepts > **Core architecture**

Efficient for same-machine communication

Simple process management

2. Remote communication

Use SSE for scenarios requiring HTTP compatibility

Consider security implications including authentication and authorization

Message handling

1. Request processing

Validate inputs thoroughly

Use type-safe schemas

Handle errors gracefully

Implement timeouts

2. Progress reporting

Use progress tokens for long operations

Report progress incrementally

Include total progress when known

3. Error management

Use appropriate error codes

Include helpful error messages

Clean up resources on errors

Security considerations



1. Transport security

Model Context Protocol

Use TLS for remote connections

Concepts > **Core architecture**

- Validate connection origins

Implement authentication when needed

2. Message validation

Validate all incoming messages

Sanitize inputs

Check message size limits

Verify JSON-RPC format

3. Resource protection

Implement access controls

Validate resource paths

Monitor resource usage

Rate limit requests

4. Error handling

Don't leak sensitive information

Log security-relevant errors

Implement proper cleanup

Handle DoS scenarios

Debugging and monitoring

1. Logging

Log protocol events

Track message flow

Monitor performance

Record errors



2. Diagnostics

Model Context Protocol

Implement health checks

Concepts > **Core architecture**

- Monitor connection state

Track resource usage

Profile performance

3. Testing

Test different transports

Verify error handling

Check edge cases

Load test servers

Was this page helpful?

 Yes

 No

[< Inspector](#)

[Resources >](#)
